Dataset.Anonymisation

December 14, 2022

1 Anonymisation of 'customer_information.csv' dataset for use by Imperial researchers and government (and calculation of K-anonymity)

Anonymisation is the practice of removing identifying information from data in order to protect individuals' privacy. By anonymising data, we can ensure that sensitive information is kept secure. In this project, we aim to create an anonymised dataset by removing personally identifiable information from the original dataset whilst attempting to retain its utility and insights by using a combination of techniques such as pseudonymous identification and data perturbation.

1.1 Helper functions

The following helper functions are needed:

```python
# Helper functions

# The following variable countries were hard-coded to fix unmatched territory
 ↪errors
northern_countries = ["Svalbard & Jan Mayen Islands"]
southern_countries = ["Micronesia"]

# Parse country into shortform
def parse_country(country_name):
    return coco.convert(country_name, to='name_short', include_obsolete=True)

# Convert country of birth into Hemisphere (Northern or Southern) based on
 ↪latitude coordinates
def country_to_hemisphere(country_name):
    try:
        if country_name in southern_countries:
            return "Southern Hemisphere"
        elif country_name in northern_countries:
            return "Northern Hemisphere"
        else:
            return ("Southern" if Nominatim(user_agent="CDM").
 ↪geocode(parse_country(country_name)).latitude < 0 else "Northern") + "␣
 ↪Hemisphere"
```

```python
        except Exception as e:
            print(e)
            return "Error"

# SHA hash function using a key and salt
def hash(to_hash, key):
    salt = os.urandom(16)
    h = hashlib.sha256()
    h.update(key)
    h.update(salt)
    h.update(to_hash.encode())
    return to_hash, h.hexdigest(), salt.hex()

# To encrypt and save as encrypted file; specify file to encrypt, encrypted
 ↪file destination, and destination key location
def encrypt(to_encrypt, file_destination, key_location):
    key = Fernet.generate_key() # AES in CBC mode with a 128-bit key for
 ↪encryption
    fernet = Fernet(key)

    with open(key_location, 'wb') as f:
        f.write(key)

    with open(to_encrypt, 'rb') as f:
        plaintext = f.read()

    encrypted = fernet.encrypt(plaintext)
    with open(file_destination, 'wb') as e:
        e.write(encrypted)
```

1.2 Loading required data and creating the anonymised dataframe

The original dataset is loaded, an empty output Dataframe is initialised, and a postcode reference table is loaded

```python
[3]: # Read in data to be anonymised
original_data = pd.read_csv("Data/customer_information.csv")

# Reading in postcode_region.csv to map given postcode to countries in the UK -
 ↪'England' and 'Other'(includes Wales, Scotland, Northern Ireland)
postcode_dictionary = pd.read_csv('Data/postcode_region.csv')

# Create anon_data variable as initial data with unneeded direct identifiers
 ↪dropped
anon_data = pd.DataFrame()

postcode_dictionary.head()
```

```
[3]:    Postcode    Region
    0         AB     Other
    1         AL   England
    2          B   England
    3         BA   England
    4         BB   England
```

1.3 Adding variables to the anonymised dataset

Assigning gender and case-control status as given

```
[4]: # Assign gender
     anon_data['Gender'] = original_data['gender']

     # Assign case-control status
     anon_data['CC.Status'] = original_data['cc_status']

     anon_data.head()
```

```
[4]:    Gender  CC.Status
    0       F          0
    1       M          0
    2       F          0
    3       F          0
    4       F          0
```

2 Anonymisation

2.1 Pseudonymisation with hashed Sample ID

Next, a unique Sample ID is created from the National Insurance Number to link the anonymised data with the reference data containing sensitive information.

```
[5]: # Clean NIN formatting and assign Sample ID as a hashed form of the NIN
     key = os.urandom(16)
     original_data["national_insurance_number"], anon_data['Sample.ID'], salts =␣
      ↪zip(*original_data["national_insurance_number"].apply(
         lambda x: hash(re.sub(r'(.{2})(?!$)','\\1 ', x.replace(' ', '') ), key)))

     anon_data.head()
```

```
[5]:    Gender  CC.Status                                       Sample.ID
    0       F          0   717341b0d3455426af5043290ec12b0681175d0b2da21a…
    1       M          0   ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
    2       F          0   dd3daf05200664054f9162a61dd76d64802bff56437e20…
    3       F          0   a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
    4       F          0   74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…
```

```
[6]: # Create a reference table between NIN and respective hashed NIN
     reference_table = pd.DataFrame()
     reference_table['Hashed.NIN'] = anon_data['Sample.ID']
     reference_table['Salt'] = salts
     reference_table['Key'] = key.hex()
     reference_table['NIN'] = original_data['national_insurance_number']

     reference_table.head()
```

```
[6]:                                        Hashed.NIN  \
     0   717341b0d3455426af5043290ec12b0681175d0b2da21a…
     1   ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
     2   dd3daf05200664054f9162a61dd76d64802bff56437e20…
     3   a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
     4   74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…

                                      Salt                                Key  \
     0   ae064c90392c2d5c545da74d82c2ae52   ae823853bfce7c66904c0de363de9b2f
     1   af9e92525c444b408b30a920c3e3fd1a   ae823853bfce7c66904c0de363de9b2f
     2   b735f1a9a61d9716d1dbed1db632d9c4   ae823853bfce7c66904c0de363de9b2f
     3   dcb27a563ea605225fc2bccdb4e7c6ea   ae823853bfce7c66904c0de363de9b2f
     4   8907a93f3a4de5bdf6dc3fd1d6de8337   ae823853bfce7c66904c0de363de9b2f

                 NIN
     0   ZZ 19 48 92 T
     1   ZZ 75 35 13 T
     2   ZZ 94 71 96 T
     3   ZZ 39 69 47 T
     4   ZZ 30 98 91 T
```

2.2 Banding

2.2.1 Date of birth and education level

```
[7]: # Birth years are extracted
     birthyears = pd.DatetimeIndex(original_data['birthdate']).year

     # Banding the birth years into 20-year intervals
     anon_data['Birthyear'] = pd.cut(birthyears, np.arange(birthyears.min(),
      ↪birthyears.max()+20, 20), right=False)

     anon_data.head()
```

```
[7]:    Gender  CC.Status                                    Sample.ID  \
     0       F          0   717341b0d3455426af5043290ec12b0681175d0b2da21a…
     1       M          0   ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
     2       F          0   dd3daf05200664054f9162a61dd76d64802bff56437e20…
```

```
3       F           0   a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
4       F           0   74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…

        Birthyear
0   [1975, 1995)
1   [1995, 2015)
2   [1975, 1995)
3   [1995, 2015)
4   [1955, 1975)
```

2.2.2 Full postcode to countries within the UK

```
[8]: # Assign UK country derived from postcode by comparing to the␣
     ↪postcode_dictionary reference table
     anon_data['Postcode'] = original_data['postcode'].apply(lambda x: re.
     ↪search('[a-zA-Z]*', x).group(0))
     anon_data = pd.merge(anon_data, postcode_dictionary, on='Postcode', how='left')
     anon_data = anon_data.rename(columns={'Region': 'UK.Country'})

     anon_data.head()
```

```
[8]:    Gender   CC.Status                                       Sample.ID  \
     0       F           0   717341b0d3455426af5043290ec12b0681175d0b2da21a…
     1       M           0   ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
     2       F           0   dd3daf05200664054f9162a61dd76d64802bff56437e20…
     3       F           0   a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
     4       F           0   74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…

          Birthyear Postcode UK.Country
     0   [1975, 1995)      LS    England
     1   [1995, 2015)       M    England
     2   [1975, 1995)      SO    England
     3   [1995, 2015)       B    England
     4   [1955, 1975)      TQ    England
```

2.2.3 Education level and country of birth

```
[9]: # Assign education level as banded education level
     anon_data['Education.Level'] = original_data['education_level'].map(lambda x:␣
     ↪"Higher" if x in ["bachelor", "masters", "phD"] else "BasicOther")

     # Assign hemisphere of birth depending on country of birth
     anon_data['Location.of.Birth'] = original_data['country_of_birth'].apply(lambda␣
     ↪x: country_to_hemisphere(x))

     anon_data.head()
```

```
[9]:    Gender  CC.Status                                             Sample.ID  \
     0      F          0  717341b0d3455426af5043290ec12b0681175d0b2da21a…
     1      M          0  ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
     2      F          0  dd3daf05200664054f9162a61dd76d64802bff56437e20…
     3      F          0  a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
     4      F          0  74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…

           Birthyear Postcode UK.Country Education.Level    Location.of.Birth
     0  [1975, 1995)       LS    England          Higher  Northern Hemisphere
     1  [1995, 2015)        M    England      BasicOther  Northern Hemisphere
     2  [1975, 1995)       SO    England          Higher  Northern Hemisphere
     3  [1995, 2015)        B    England      BasicOther  Northern Hemisphere
     4  [1955, 1975)       TQ    England      BasicOther  Southern Hemisphere
```

2.3 Data perturbation (adding Gaussian noise)

```python
[10]: # Add gaussian noise to weight, height, countries visited, average number of␣
       ↪drinks in alcohol units per week and average cigrettes smoked per week.
      weight_noise = np.random.normal(0,1,1000)*5
      anon_data['Weight'] = round(original_data['weight']+weight_noise, 1)


      height_noise = np.random.normal(0,1,1000)/5
      anon_data['Height'] = round(original_data['height']+height_noise, 2)


      countries_noise = np.random.normal(0,1,1000)*5
      anon_data['Countries.Visited'] =␣
       ↪round(original_data['n_countries_visited']+countries_noise)


      alcohol_noise = np.random.normal(0,1,1000)
      anon_data['Avg.Alcohol'] =␣
       ↪round(original_data['avg_n_drinks_per_week']+alcohol_noise, 1)


      smoking_noise = np.random.normal(0,1,1000)*20
      anon_data['Avg.Cigarettes'] =␣
       ↪round(original_data['avg_n_cigret_per_week']+smoking_noise)


      anon_data.head()
```

```
[10]:    Gender  CC.Status                                             Sample.ID  \
     0      F          0  717341b0d3455426af5043290ec12b0681175d0b2da21a…
     1      M          0  ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…
     2      F          0  dd3daf05200664054f9162a61dd76d64802bff56437e20…
     3      F          0  a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…
     4      F          0  74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…

           Birthyear Postcode UK.Country Education.Level    Location.of.Birth  \
```

```
0  [1975, 1995)         LS     England              Higher  Northern Hemisphere
1  [1995, 2015)          M     England         BasicOther  Northern Hemisphere
2  [1975, 1995)         SO     England              Higher  Northern Hemisphere
3  [1995, 2015)          B     England         BasicOther  Northern Hemisphere
4  [1955, 1975)         TQ     England         BasicOther  Southern Hemisphere

    Weight  Height  Countries.Visited  Avg.Alcohol  Avg.Cigarettes
0     73.5    1.58               39.0          4.3           208.0
1     66.9    1.52               43.0          1.5            50.0
2     92.0    1.84                7.0          6.3            61.0
3     58.9    1.84               33.0          4.4           261.0
4    101.9    1.74               29.0          4.9           348.0
```

3 Calculating K-anonymity using quasi-identifiers

The following code groups the quasi-identifiers specified, calculates the k-value, and returns a count of the "unique" rows.

```python
[66]: # Checking counts of quasi-identifiers permutations
df_count = anon_data.groupby(['Gender', 'Birthyear', 'Location.of.Birth',
                              'UK.Country', 'Education.Level']).size().
  ↪reset_index(name = 'Count')
counts_exceeding_0 = df_count.sort_values("Count")[df_count.
  ↪sort_values("Count")['Count'] > 0].reset_index()

# Calculate the k-value and the number of unique quasi-identifier permutations
print("The anonymised dataset is " + str(counts_exceeding_0['Count'][0]) +␣
  ↪"-anonymous; there are " +
      str(len(df_count[df_count['Count']==1])) + " unique quasi-identifier␣
  ↪permutations.\n")

# Printing the final grouped output
print("Least-frequent quasi-identifier permutations, in ascending order:")
counts_exceeding_0.head()
```

The anonymised dataset is 2-anonymous; there are 0 unique quasi-identifier permutations.

Least-frequent quasi-identifier permutations, in ascending order:

```
[66]:    index Gender      Birthyear     Location.of.Birth UK.Country Education.Level  \
    0      47      M  [1995, 2015)  Southern Hemisphere      Other          Higher
    1      31      M  [1955, 1975)  Southern Hemisphere      Other          Higher
    2      46      M  [1995, 2015)  Southern Hemisphere      Other      BasicOther
    3      19      F  [1995, 2015)  Northern Hemisphere      Other          Higher
    4      39      M  [1975, 1995)  Southern Hemisphere      Other          Higher
```

```
     Count
0        2
1        2
2        2
3        2
4        3
```

4   Viewing and saving the anonymised dataset

```
[47]: # Re-order columns
      anon_data = anon_data[['Sample.ID', 'Gender', 'Birthyear', 'Location.of.Birth',␣
       ↪'UK.Country', 'Weight',
                             'Height', 'Education.Level', 'Avg.Alcohol', 'Avg.
       ↪Cigarettes', 'Countries.Visited', 'CC.Status']]

      # View the anonymised dataset
      anon_data.head()
```

```
[47]:                                             Sample.ID Gender    Birthyear  \
      0  717341b0d3455426af5043290ec12b0681175d0b2da21a…      F  [1975, 1995)
      1  ff878150b526edf6a15e3bcc58e6a30bd2852f5dbef10c…      M  [1995, 2015)
      2  dd3daf05200664054f9162a61dd76d64802bff56437e20…      F  [1975, 1995)
      3  a17d9326ef0ada45cfbc51fdf59c70129813f7c4bf2421…      F  [1995, 2015)
      4  74e8b816673182819be3ba63b609ff22fb9b8f8a28426f…      F  [1955, 1975)

           Location.of.Birth UK.Country  Weight  Height Education.Level  \
      0  Northern Hemisphere    England    73.5    1.58          Higher
      1  Northern Hemisphere    England    66.9    1.52      BasicOther
      2  Northern Hemisphere    England    92.0    1.84          Higher
      3  Northern Hemisphere    England    58.9    1.84      BasicOther
      4  Southern Hemisphere    England   101.9    1.74      BasicOther

         Avg.Alcohol  Avg.Cigarettes  Countries.Visited  CC.Status
      0          4.3           208.0               39.0          0
      1          1.5            50.0               43.0          0
      2          6.3            61.0                7.0          0
      3          4.4           261.0               33.0          0
      4          4.9           348.0               29.0          0
```

4.1   Creating CSV files for the anonymised data and the reference table

```
[48]: # Output the files into .csv format
      output_name = "anon_dataset"
      anon_data.to_csv(output_name + ".csv", sep=",", index=None)
```

```python
reference_table.to_csv("reference_table.csv", sep=",", index=None)
```

4.2   Encrypting the dataset

```python
[49]: # Encrypt csv and delete original file
      encrypt(output_name + ".csv", output_name + "_encrypted.csv", "key.key")
      os.remove(output_name + ".csv")
```

```python
[50]: pip freeze > requirements.txt
```

Note: you may need to restart the kernel to use updated packages.